Session 1.4: Tabular Data Wrangling Part 2: Advanced Filtering, Multi-Segment GroupBy, and Financial Auditing

THE BUSINESS SITUATION

The Quarterly Budget Review

- Each division gets a spending budget every quarter
- Finance compares the budget with what was actually spent
- Our case: 4 divisions × 4 quarters = 16 rows in one CSV file
- The finance director asks: who overspent, by how much, and who is responsible?

Today, Python answers this for the whole year at once.

Today's flow: 60 min lecture, then 120 min hands-on lab

EXPECTED OUTCOMES

By the end of today, you will be able to:

1. ***"Which quarters went over budget?"***
   Filter rows with a Boolean mask and `&`
   Boolean mask: a `True`/`False` value for every row
2. ***"How much did each division overspend this year?"***
   Summarize groups with `groupby()` and `.agg()`
   `groupby()`: split rows into groups, like an Excel pivot table (a summary with one row per group)
3. ***"Who is responsible?"***
   Join tables with `pd.merge()` and fix a `KeyError`
   `pd.merge()`: join two tables on a shared column, like Excel's `VLOOKUP` lookup formula

Session Map: Follow the Numbers

Today's question: *Who overspent, by how much, and who is responsible?*

1 Why it matters
Lecture, part 1

1. Session Purpose
2. Learning Outcomes
3. Business Context: Budget vs Actual
4. Data Basics: From Excel Tools You Know to pandas

→

2 Learn the tools
Lecture, part 2

5. Lab Practice Dataset
6. Keywords
7. Core Foundations
8. IDE Workflow
9. Code Examples + Quick Checks
10. Official References

→

3 Practice
Lab

11. Exercise 1: Guided - Quarterly Division Audit
12. Exercise 2: Applied - Annual Cost Center Overruns
13. Exercise 3: Advanced - VP Exception Report

Budget vs Actual: How Finance Spots Overspending

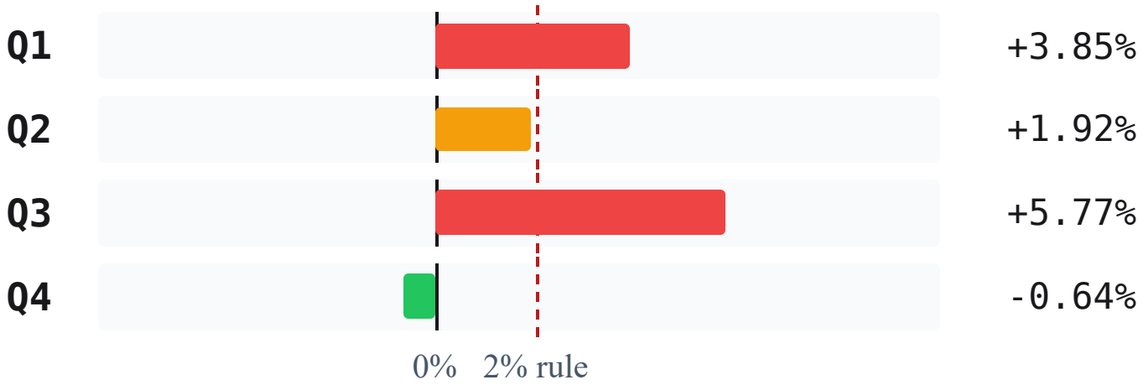
FOUR TERMS, ONE EXAMPLE

Consumer Division, Q1

Budget	planned spending	`780,000`
Actual	real spending	`810,000`
Variance	Actual - Budget	`+30,000`
Variance %	Variance / Budget × 100	`+3.85%`

- Positive variance = **overspent**; negative = under budget
- Cost center: the accounting code spending is charged to; each has a VP (Vice President)
- Company rule in our data: flag a quarter when variance is **more than 2%** (column `Audit_Status`)

CONSUMER DIVISION: VARIANCE % BY QUARTER



Red: flagged (over 2%) **Amber**: over, within rule

Green: under budget

Today's job: find these overruns for all 4 divisions, add up the year, and name the person responsible.

From Excel Tools You Know to pandas

EXCEL

Filter

Click the column arrow, tick the boxes

→

PANDAS

Boolean mask: `df[mask]`

The rule is written in code, so it repeats exactly next quarter

EXCEL

Pivot Table

A summary by group: `Division_Name` in Rows, spending in Values = 4 rows instead of 16

→

PANDAS

`df.groupby().agg()`

One line; run it again when new data arrives

EXCEL

VLOOKUP

A formula that looks up a code in another table and returns column **number** N, such as the 2nd

→

PANDAS

`pd.merge()`

Matches rows by column **name**, such as `Cost_Center_Code`

One File and One Small Table

THE FILE: `session_01_04_dataset.csv` (16 ROWS)

`Transaction_ID`	Record code, such as `TXN-5001`
`Division_Name`	Cloud, Consumer, Enterprise, Logistics
`Cost_Center_Code`	Accounting code, such as `CC-202`
`Fiscal_Quarter`	Q1 to Q4
`Budget_Expense_USD`	Planned spending
`Actual_Expense_USD`	Real spending
`Audit_Status`	`Variance_Flag` if more than 2% over budget

THE SMALL TABLE: `vp_info`

- Built in code in Demo 3 with `pd.DataFrame({...})`
- 4 rows: one VP (Vice President) per cost center
- VP names are fictional
- Used in: Demo 3, Exercises 2-3

Each CSV row is one division's spending in one quarter: 4 divisions × 4 quarters = 16 rows.

Six Terms You Need Today

FILTER

Boolean Mask

What: one `True` or `False` per row

Why: picks exactly the rows that meet a rule

COMBINE RULES

`&` `|` `~`

What: "and", "or", "not" for masks

Why: Python's `and` / `or` fail on whole columns

PATTERN

Split-Apply-Combine

What: split rows into groups, calculate each group, combine the results

Why: the idea behind every pivot table

SUMMARY

`groupby()` + `.agg()`

What: `groupby()` splits; `.agg()` sums, counts, or averages each group

Why: yearly totals per division in one step

JOIN

`pd.merge()`

What: joins two tables row by row

Why: adds facts from another table, such as the VP's name, with no copy-paste

MATCH

Key Column

What: a column both tables share, such as `Cost_Center_Code`

Why: tells pandas which rows belong together

How the Six Terms Fit Together

STEP 1: BOOLEAN MASK (CONSUMER ROWS)

Fiscal_Quarter	Over budget?
Q1	True
Q2	True
Q3	True
Q4	False

Blue = `True`, kept by `df[mask]`

STEP 2: SPLIT, APPLY, COMBINE

Split: 16 rows

- `Cloud × 4`
- `Consumer × 4`
- `Enterprise × 4`
- `Logistics × 4`

→ sum

Division_Name	Variance_USD
Cloud	-15000.0
Consumer	85000.0
Enterprise	0.0
Logistics	28000.0

Combine: one row per division for the whole year

STEP 3: MERGE ON A KEY COLUMN

Cost_Center_Code	Variance_USD
CC-202	85000.0
CC-404	28000.0

+

Cost_Center_Code	VP_Lead
CC-202	Marcus Vance
CC-404	David Kim

=

Cost_Center_Code	Variance_USD	VP_Lead
CC-202	85000.0	Marcus Vance
CC-404	28000.0	David Kim

Yellow = the key column both tables share; green = the new information. VP names are fictional.

How Filtering, Grouping, and Joining Work

7.1

`&`, Not `and`

- A rule gives one `True` / `False` per row
- `and` raises `ValueError`; no brackets raises `TypeError`

Remember: combine with `&` `|` `~`, each rule in `( )`

7.2

Split, Apply, Combine

- Consumer: 4 quarters, summed to `3205000.0`
- 16 rows become 4, one per division

Remember: `groupby()` splits, `.agg()` applies

7.3

Match by Name, Not Position

- `VLOOKUP` returns "column 2": breaks if a column is inserted
- `pd.merge()` matches the key column by name

Remember: the key needs the same name in both tables

7.4

`how="left"` vs `"inner"`

- `left`: an unmatched row stays, with `NaN`
- `inner`: the row disappears from the report

Remember: use `left`, then look for `NaN`

Filtering, Grouping, and Joining, Seen Up Close

1. `and` FAILS, `&` WORKS

```
is_consumer = df["Division_Name"] == "Consumer"
is_over = df["Actual_Expense_USD"] > df["Budget_Expense_USD"]
```

```
is_consumer and is_over
```

`ValueError: The truth value of a Series is ambiguous`

```
is_consumer & is_over
```

One `True` / `False` per row

2. APPLY: CONSUMER ACTUAL SPENDING

Actual_Expense_USD
810000.0
795000.0
825000.0
775000.0

→ `"sum"`

Total
3205000.0

Same step for all 4 divisions

3. VLOOKUP BREAKS WHEN A COLUMN IS INSERTED

Before

Code	VP
CC-202	Marcus Vance

After inserting Region

Code	Region	VP
CC-202	North	Marcus Vance

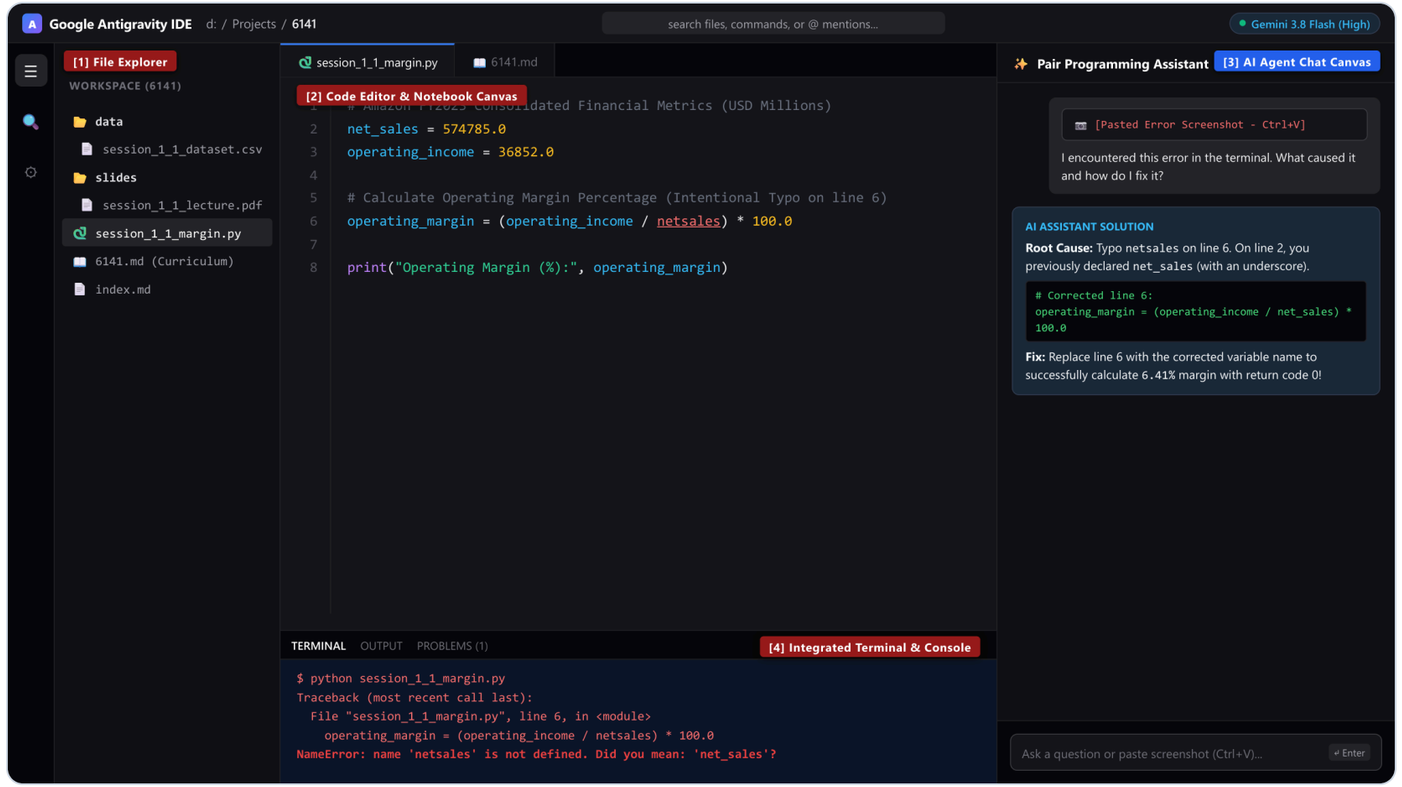
"2nd column" now returns "North": wrong, and easy to miss

4. `how="left"` KEEPS UNMATCHED ROWS

Cost_Center_Code	VP_Lead
CC-202	Marcus Vance
CC-404	David Kim
CC-505	*NaN*

A code with no VP (example `CC-505`) stays, with `NaN`

Build, Test, and Diagnose in the IDE



Example: Terminal `[4]` shows a `NameError`; AI Chat `[3]` explains the fix

[1] File Explorer

Name it `YourName_S14_demo_1.py`
Example: `Reacher_S14_demo_1.py`
Check the dataset is in `data/`

[2] Code Editor

Write a few lines, run, repeat
Small steps keep errors small

[3] AI Agent Chat

Error? Capture the traceback (Python's error report) with `Win + Shift + S`
Paste with `Ctrl + V`, ask why

[4] Terminal

Run `python` + your script name
Return code `0` = no errors

Find Overspending with a Boolean Mask

YourName_S14_demo_1.py — Input → Process → Output

```
# 1. Input: load the quarterly budget data
import pandas as pd
df = pd.read_csv("data/session_01_04_dataset.csv")

# 2. Process: Consumer quarters that spent more than budget
mask = (
    (df["Division_Name"] == "Consumer")
    & (df["Actual_Expense_USD"] > df["Budget_Expense_USD"])
)
overruns = df[mask]

# 3. Output: count and list them
print("Consumer overrun quarters:", len(overruns))
print(overruns[["Fiscal_Quarter", "Budget_Expense_USD",
                "Actual_Expense_USD"]])
```

1. INPUT

Same load step as Session 1.3

`pd.read_csv()` reads the 16 rows into `df`

2. PROCESS

`mask` and `df[mask]`

- `mask`: one `True` / `False` per row
- `&`: both rules must be `True`
- `df[mask]` keeps only the `True` rows

3. OUTPUT

`len()` and a short table

The count, then the matching quarters

Flip a Rule with `~`

NEW TOOL

`~` (not)

What: turns every `True` into `False`, and every `False` into `True`

Why: find the opposite group, such as quarters that stayed within budget

Fiscal_Quarter	Over budget?	After `~`
Q1	True	False
Q2	True	False
Q3	True	False
Q4	False	True

Consumer division; blue = `True` (kept)

YOUR TASK

Change One Line

In `YourName_S14_demo_1.py`, part **2. Process**: find Consumer quarters that did **not** go over budget.

Now

```
    & (df["Actual_Expense_USD"] > df["Budget_Expense_USD"])
```

Your change: fill in the yellow blank

```
    & ___(df["Actual_Expense_USD"] > df["Budget_Expense_USD"])
```

Check: the script prints `Consumer overrun quarters: 1`, and the table shows `Q4` (budget `780000.0`, actual `775000.0`). The label still says "overrun", since only the rule changed.

Total Each Division with `groupby()`

YourName_S14_demo_2.py — Input → Process → Output

```python
# 1. Input: load the quarterly budget data
import pandas as pd
df = pd.read_csv("data/session_01_04_dataset.csv")

# 2. Process: yearly totals for each division
div_summary = df.groupby("Division_Name").agg(
    Total_Budget_USD=("Budget_Expense_USD", "sum"),
    Total_Actual_USD=("Actual_Expense_USD", "sum"),
    Quarters=("Transaction_ID", "count"),
)
div_summary["Variance_USD"] = (
    div_summary["Total_Actual_USD"]
    - div_summary["Total_Budget_USD"]
)

# 3. Output: one row per division
print(div_summary[["Total_Actual_USD", "Variance_USD"]])
```

1. INPUT

Same load step

Every script loads its own data

2. PROCESS

`groupby()` + `.agg()`

- `groupby("Division_Name")`: 4 groups
- Inside `.agg()`: `New_Name=("Column", "sum")`
- Then one subtraction per division

3. OUTPUT

One row per division

4 rows instead of 16

Rank Results with `.sort_values()`

NEW TOOL

`.sort_values()`

What: reorders rows by one column; `ascending=False` puts the largest first

Why: the director reads the biggest overspender first

Division_Name		Division_Name
Cloud		?
Consumer	→	?
Enterprise		?
Logistics		?

Now: A to Z. After sorting: largest variance on top

YOUR TASK

Add One Line

In `YourName_S14_demo_2.py`, add one line at the end of **2. Process**: largest `Variance_USD` first.

Now

```
# (new line goes here)
```

Your change: fill in the yellow blank

```
div_summary = div_summary.____________(
    "Variance_USD", ascending=False
)
```

Check: the first row is `Consumer` (`85000.0`), then `Logistics` (`28000.0`).

Add VP Names with `pd.merge()`

YourName_S14_demo_3.py — Input → Process → Output

```
# 1. Input: budget data and a VP list (fictional names)
import pandas as pd
df = pd.read_csv("data/session_01_04_dataset.csv")
vp_info = pd.DataFrame({
    "Cost_Center_Code": ["CC-101", "CC-202", "CC-303", "CC-404"],
    "VP_Lead": ["Sarah Chen", "Marcus Vance",
                "Elena Rostova", "David Kim"],
})

# 2. Process: add VP names, then flag spending over budget
ledger = pd.merge(df, vp_info, on="Cost_Center_Code", how="left")
ledger["Over_Budget"] = (
    ledger["Actual_Expense_USD"] > ledger["Budget_Expense_USD"]
)

# 3. Output: show the first three rows
print(ledger[["Transaction_ID", "VP_Lead", "Over_Budget"]].head(3))
```

1. INPUT

Two tables

- The CSV file
- `pd.DataFrame({...})` builds a small table in code: column name, then its list of values

2. PROCESS

`pd.merge()` + a flag

- `on=` names the key column
- `how="left"` keeps all 16 rows
- `>` makes a `True` / `False` column

3. OUTPUT

`.head(3)`

The first rows, now with `VP_Lead`

Count Categories with `.value_counts()`

NEW TOOL

`.value_counts()`

What: counts how often each value appears in a column

Why: answers "how many quarters were over budget?" in one line

Over_Budget
False
True
False
... 16 rows

→

Value	Count
False	?
True	?

One count per distinct value

YOUR TASK

Add One Line

In `YourName_S14_demo_3.py`, add one line to **3. Output**: count the `Over_Budget` values.

Now

```
# (new line goes here)
```

Your change: fill in the yellow blank

```
print(ledger["Over_Budget"].[redacted]())
```

Check: the Terminal shows `False 8` and `True 8`: half of the 16 quarters went over budget.

Where to Look Things Up

PANDAS

GroupBy: Split-Apply-Combine

Everything `groupby()` and `.agg()` can do. pandas.pydata.org/docs/user_guide/groupby.html

PANDAS

Merge, Join, Concatenate and Compare

All the options of `pd.merge()`, including `how=`. pandas.pydata.org/docs/user_guide/merging.html

PYTHON STYLE GUIDE

PEP 8: Line Breaks Around Operators

Why we start a new line with `&` or `-` in long rules. peps.python.org/pep-0008

Exercise 1 and Exercise 2

EXERCISE 1: GUIDED PRACTICE

Quarterly Division Audit

- Create `YourName_S14_audit.py`
- Add `Variance_USD` and `Variance_Pct` columns
- Group by two columns: `["Division_Name", "Fiscal_Quarter"]`
- Sum actual spending and variance with `.agg()`
- Run it in Terminal `[4]`

EXERCISE 2: APPLIED PRACTICE

Annual Cost Center Overruns

- In `YourName_S14_practice.py`, group by `Cost_Center_Code` and compute the yearly variance %
- Keep cost centers more than `1.0` % over budget
- Merge with `vp_info`, but type `on="Cost_Center"` on purpose; read the `KeyError` and ask AI Chat `[3]`
- Fix the key and print cost center, VP, and overrun

Success check: 2 cost centers: `CC-202` (Marcus Vance, `85000.0`, 2.72%) and `CC-404` (David Kim, `28000.0`, 1.15%); return code `0`.

VP Exception Report: Who Must Explain?

WHAT TO BUILD: `YourName_S14_challenge.py`

1. `Variance_USD`, `Variance_Pct` (2 decimals); `My_Flag` = `Variance_Pct > 2.0`; check it agrees with `Audit_Status`
2. Flagged rows only, by cost center: count, overrun sum, worst %
3. All rows, by cost center: `Year_Variance_Pct`
4. Merge both into `vp_info` (`how="left"`)
5. `Escalate` = 2+ flagged quarters **and** year variance `> 1.0`
6. Sort, print, save `data/vp_exception_report.csv`

NEW TOOLS

- `.reset_index()`: group names back into a normal column
- `.fillna(0)`: empty values become `0`
- `"max"` in `.agg()`: the largest value
- Two sort keys:

```
.sort_values(
    ["Flagged_Overrun_USD",
     "Worst_Variance_Pct"],
    ascending=False,
)
```

Self-check: 0 disagreements; Escalate only for Marcus Vance (2.72%) and David Kim (1.15%). Full steps: section 13 of the course page.